

```
In [5]: ▶ import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
In [6]: ▶ df=pd.read_csv("creditcard.csv")
```

```
In [7]: ▶ df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
Time          284807 non-null float64
V1            284807 non-null float64
V2            284807 non-null float64
V3            284807 non-null float64
V4            284807 non-null float64
V5            284807 non-null float64
V6            284807 non-null float64
V7            284807 non-null float64
V8            284807 non-null float64
V9            284807 non-null float64
V10           284807 non-null float64
V11           284807 non-null float64
V12           284807 non-null float64
V13           284807 non-null float64
V14           284807 non-null float64
V15           284807 non-null float64
V16           284807 non-null float64
V17           284807 non-null float64
V18           284807 non-null float64
V19           284807 non-null float64
V20           284807 non-null float64
V21           284807 non-null float64
V22           284807 non-null float64
V23           284807 non-null float64
V24           284807 non-null float64
V25           284807 non-null float64
V26           284807 non-null float64
V27           284807 non-null float64
V28           284807 non-null float64
Amount        284807 non-null float64
Class         284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [8]: df.head()
```


Out[8]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533

5 rows × 31 columns

```
In [9]: sns.countplot(x='Class',data=df)
```


Out[9]: <matplotlib.axes._subplots.AxesSubplot at 0x29db574a668>



```
In [10]: #sns.pairplot(df.sample(1000))
```

```
In [11]: sns.heatmap(df.corr())
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x29db63553c8>
```




```
In [12]: df.drop('Time',axis=1,inplace=True)
```

```
In [13]: df.head(1)
```

```
Out[13]:
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787

1 rows × 30 columns



```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [15]: X=df.drop('Class',axis=1)
y=df['Class']
```

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, ran
```

```
In [17]: from sklearn.linear_model import LinearRegression
```

```
In [18]: lm = LinearRegression()
```


```
In [19]: lm.fit(X_train,y_train)
```

```
Out[19]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
In [20]: predictions = lm.predict(X_test)
```

```
In [21]: plt.scatter(y_test,predictions)
```

```
Out[21]: <matplotlib.collections.PathCollection at 0x29db68c76d8>
```



```
In [22]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
```

```
In [23]: ▶ print(classification_report(y_test,predictions))
```

```
-----  
ValueError                                Traceback (most recent call last)  
<ipython-input-23-49b4b26ceb45> in <module>()  
----> 1 print(classification_report(y_test,predictions))  
  
F:\anaconda\lib\site-packages\sklearn\metrics\_classification.py in classif  
ication_report(y_true, y_pred, labels, target_names, sample_weight, digits,  
output_dict, zero_division)  
    1969     """  
    1970  
-> 1971     y_type, y_true, y_pred = _check_targets(y_true, y_pred)  
    1972  
    1973     labels_given = True  
  
F:\anaconda\lib\site-packages\sklearn\metrics\_classification.py in _check_  
targets(y_true, y_pred)  
     88     if len(y_type) > 1:  
     89         raise ValueError("Classification metrics can't handle a mix  
of {0} "  
---> 90                                     "and {1} targets".format(type_true, type_p  
red))  
     91  
     92     # We can't have more than one value on y_type => The set is no  
more needed  
  
ValueError: Classification metrics can't handle a mix of binary and continu  
ous targets
```

```
In [24]: ▶ from sklearn.linear_model import LogisticRegression
```

```
In [25]: ► logmodel = LogisticRegression()  
logmodel.fit(X_train,y_train)
```

```
F:\anaconda\lib\site-packages\sklearn\linear_model\_logistic.py:940: Conver  
genceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:  
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessing.html)  
Please also refer to the documentation for alternative solver options:  
    https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression)  
    extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)
```

```
Out[25]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                             intercept_scaling=1, l1_ratio=None, max_iter=100,  
                             multi_class='auto', n_jobs=None, penalty='l2',  
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,  
                             warm_start=False)
```

```
In [26]: ► predictions = logmodel.predict(X_test)
```

```
In [27]: ► print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71077
1	0.88	0.61	0.72	125
accuracy			1.00	71202
macro avg	0.94	0.80	0.86	71202
weighted avg	1.00	1.00	1.00	71202

```
In [28]: ► print(confusion_matrix(y_test,predictions))
```

```
[[71067  10]  
 [  49   76]]
```

```
In [29]: ► from sklearn.neighbors import KNeighborsClassifier
```

```
In [30]: ► knn = KNeighborsClassifier(n_neighbors=1)
```

```
In [31]:  ➤ knn.fit(X_train,y_train)
```

```
Out[31]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',  
                               metric_params=None, n_jobs=None, n_neighbors=1, p=2,  
                               weights='uniform')
```

```
In [32]:  ➤ pred = knn.predict(X_test)
```

```
In [33]:  ➤ print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71077
1	0.88	0.73	0.80	125
accuracy			1.00	71202
macro avg	0.94	0.86	0.90	71202
weighted avg	1.00	1.00	1.00	71202

```
In [34]:  ➤ print(confusion_matrix(y_test,pred))
```

```
[[71065   12]  
 [   34   91]]
```

```
In [ ]:  ➤ error_rate = []
```

```
# Will take some time  
for i in range(1,10):  
  
    knn = KNeighborsClassifier(n_neighbors=i)  
    knn.fit(X_train,y_train)  
    pred_i = knn.predict(X_test)  
    error_rate.append(np.mean(pred_i != y_test))
```

```
In [35]:  ➤ from sklearn.tree import DecisionTreeClassifier
```

```
In [36]:  ➤ dtree = DecisionTreeClassifier()
```

```
In [38]:  ➤ dtree.fit(X_train,y_train)
```

```
Out[38]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',  
                                 max_depth=None, max_features=None, max_leaf_nodes=None,  
                                 min_impurity_decrease=0.0, min_impurity_split=None,  
                                 min_samples_leaf=1, min_samples_split=2,  
                                 min_weight_fraction_leaf=0.0, presort='deprecated',  
                                 random_state=None, splitter='best')
```

```
In [39]:  ➤ predictions = dtree.predict(X_test)
```

```
In [40]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71077
1	0.72	0.78	0.75	125
accuracy			1.00	71202
macro avg	0.86	0.89	0.88	71202
weighted avg	1.00	1.00	1.00	71202

```
In [41]: print(confusion_matrix(y_test,predictions))
```

```
[[71039   38]
 [   27   98]]
```

```
In [42]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(n_estimators=100)
rfc.fit(X_train, y_train)
```

```
Out[42]: RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                                criterion='gini', max_depth=None, max_features='auto',
                                max_leaf_nodes=None, max_samples=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=100,
                                n_jobs=None, oob_score=False, random_state=None,
                                verbose=0, warm_start=False)
```

```
In [43]: rfc_pred = rfc.predict(X_test)
```

```
In [44]: print(confusion_matrix(y_test,rfc_pred))
```

```
[[71070    7]
 [   25  100]]
```

```
In [45]: print(classification_report(y_test,rfc_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	71077
1	0.93	0.80	0.86	125
accuracy			1.00	71202
macro avg	0.97	0.90	0.93	71202
weighted avg	1.00	1.00	1.00	71202

```
In [ ]: 
```